SquadStack – Vertical Expansion Strategy

Date: March 11, 2026
Role: Founder's Office Application
Methodology: Web Research + SquadStack Fundamental Equations + Segment Analysis
Standard: No company financials. Pure segment opportunity analysis.
Prepared By: Ayush Lahoti

Executive Summary

The Strategic Question

Should SquadStack expand into Real Estate, EdTech, and Healthcare in parallel, or prioritize a single new beachhead while continuing to deepen existing fintech and logistics business?

Short Answer

Pursue a **Real Estate-first expansion**, run **structured pilots in Healthcare in parallel**, and treat **EdTech as a longer-term, human-intensive bet** requiring careful supply-side design before full commitment.

Strategic Recommendation

Rank	Vertical	ARPU (per logo)	Margin Band	LTV/CAC Band	Y1 Realistic	Y1 Ideal
🥇 #1	Real Estate	₹40–80L	70–80%	8–12x	₹2–4Cr	₹6–10Cr
🥈 #2	Healthcare	₹15–30L	65–80%	6–10x	₹1–2Cr	₹4–6Cr
🥉 #3	EdTech	₹50–100L	60–75%	5–9x	₹1–3Cr	₹4–8Cr

Why This Ranking

Real Estate wins on outcome-pricing potential. CPL in Tier 1 markets ranges from ₹1,000–₹2,500 and cost-per-site-visit reaches ₹6,000–₹12,000. SquadStack's AI can call within seconds, filter junk leads at near-zero cost, and route high-intent leads to human agents for visit scheduling — with a clear, defensible Success Fee model.

Healthcare is the cleanest AI-only use case. Appointment booking, reminders, and FAQ handling are highly standardized. A study of 40+ Indian clinics showed 18–27% missed-call rates causing ~₹7.4L annual revenue leakage per clinic. The problem is structural and urgent. However, low ACV per account and fragmented buyers constrain near-term revenue scale.

EdTech has huge TAM but constrained margins. Support can consume 15–25% of revenue at large platforms, creating extreme urgency. But EdTech sales require emotionally intelligent human counselors — AI handles FAQs but cannot close a high-stakes enrollment with a parent. This limits AI coverage and keeps variable costs high.

Path Forward

1. **Weeks 1–12 (Pilot 1):** Run scoped pilots across all 3 verticals. Validate automation %, contact rates, and early conversion uplift.
2. **Weeks 13–24 (Pilot 2):** Double down on the top 1–2 verticals. Validate pricing power, CAC, and LTV/CAC.
3. **Week 24 (Decision Gate):** Scale, pause, or pivot each vertical based on hard unit-economic thresholds.

Part 1 – Fundamental Equation for New Verticals

1.1 Unit-Level Economics

SquadStack's core engine converts raw enterprise leads into qualified outcomes using Voice AI and human agents. The **atomic unit** is a processed lead; the **economic scaling unit** is an enterprise campaign.

Lead-level equation:

```
Profit_Lead = Revenue_Lead - Cost_Lead

Revenue_Lead = (Talk_Minutes x Price_Per_Minute)

             + (Conversion_Probability x Success_Fee)

Cost_Lead    = Routing_and_AI_Cost

             + (Human_Talk_Minutes x Agent_Payout)

             + Telecom_and_API_Cost
```

Campaign-level equation for any vertical (v):

Profit_Campaign_v

= SUM(Profit_Lead_i) + Campaign_Fixed_Revenue_v - Campaign_Fixed_Costs_v

Campaign_Fixed_Revenue_v = Setup_Fees_v + Platform_Minimums_v

Campaign_Fixed_Costs_v = Integration_Cost_v

+ Playbook_Design_v

+ Agent_Sourcing_and_Training_v

+ Ongoing_Account_Ops_v

Key economic insight: The primary cost lever is `Human_Talk_Minutes`. Every minute shifted from a human agent to AI pre-qualification directly expands margin. The primary revenue lever is `Conversion_Probability x Success_Fee` — the quality of routing determines how much of the Success Fee is actually earned.

1.2 Customer-Level Viability (LTV/CAC)

LTV_v = Annual_Profit_Account_v x Lifetime_Years_v x Retention_Rate_v

LTV/CAC_v = LTV_v / CAC_v

Interpretation thresholds:

- **LTV/CAC >= 3x** — Sustainable baseline.
- **LTV/CAC >= 8–10x** — Capital-efficient; supports aggressive scaling.
- **LTV/CAC < 3x** — Value destruction.

For any new vertical to warrant SquadStack's investment, the unit economics must clear the 8x threshold given the 10x growth mandate over 18 months.

1.3 The Six Levers (Ranked by Profit Impact)

If improved by 15–20%, the following levers drive outsized impact across all new verticals:

Rank	Lever	Mechanism	Timeline
1	AI Coverage of Talk Minutes	Reduces Human_Talk_Minutes x Agent_Payout	Immediate
2	Conversion Probability	Lifts Success_Fee yield per lead	1–3 months
3	Agent Time-to-Productivity	Shrinks Agent_Sourcing_Training per campaign	2–4 months
4	Lead Validity Rate	Protects per-lead economics from dead leads	Immediate
5	Outcome-Based Pricing Mix	Shifts revenue upside from volume to value	3–6 months
6	Vertical Mix	Reallocates capacity toward highest-margin segments	6–18 months

---

Part 2 – Segment 1: Real Estate / PropTech

2.1 Market Opportunity: TAM → SAM → SOM

TAM (India, telecalling around property sales)

Developers and PropTech platforms depend heavily on performance-marketing leads from Meta and Google and in-house or outsourced call centers to convert them into site visits. In Tier 1 markets, CPL ranges from ₹1,000–₹2,500 and cost-per-site-visit can reach ₹6,000–₹12,000 due to low connect rates and slow follow-up. Aggregated across major metros and large developers, the telecalling + lead-gen opportunity around housing sales is a multi-hundred crore annual market.

SAM (Serviceable)

Focus: Top 50–100 developers and large PropTechs with:

- Centralized lead management systems.
- High monthly digital lead volumes (500+ leads/month per project).
- Multi-city presence enabling portfolio expansion.

Average telecalling spend equivalent: ₹2–5Cr/year per large developer. Realistically, 30–40 such accounts could adopt AI+Human orchestration within 2–3 years.

SOM (Year 1 Obtainable)

- Realistic Y1: 3–5 logos → **₹2–4Cr**
- Ideal Y1: 8–10 logos (with portfolio expansions) → **₹6–10Cr**

2.2 Demand Validation – Is the Problem Real?

Core pain points verified from market data:

- **Slow first-response kills leads.** If a developer does not call within 5 minutes of a digital lead filling a form, the lead quality drops sharply. Human agents can't scale for this.
- **High cost-per-visit.** Paying ₹2,000 for a lead and ₹10,000 for a visit — when standard in-house conversion rates hover at 10–15% — creates extreme pressure.
- **Inconsistent agent quality.** High attrition in in-house calling teams means new agents are always ramping up on expensive leads.

Call-type breakdown (typical residential developer):

Call Type	Share	AI-Suitable?
Basic qualification (budget, BHK, timeline)	40–50%	Yes
Follow-up and scheduling confirmations	20–30%	Yes
Complex objection / negotiation	20–30%	Human preferred

Verdict: Strong demand. Problem is expensive, well-understood, and structurally recurring.

2.3 Indicative Unit Economics – Real Estate

Example: NCR developer with one flagship project

Assumptions:

- 3,000 new leads/month from digital campaigns.
- Baseline: 25% reachability, 10% visit conversion among reached leads.
- With SquadStack: 40% reachability, 15–18% visit conversion.

Revenue side:

- Platform minimum + usage fees: ₹20–30L/year.
- Success fee (₹750–₹1,250 per verified visit): estimated 500–700 visits/year = ₹3.75–8.75L/year.
- Indicative **ACV per project**: ₹40–60L; scales with portfolio.

Cost side:

- AI + Telecom/API costs per call cluster.
- Human payout focused only on high-intent leads.
- Campaign setup amortized over multiple months and similar playbooks.

Directional result: Real Estate campaigns can plausibly deliver **70–80% gross margin** with strong upside from outcome-based success fees.

Profit_Campaign_RealEstate

= (Visits x Success_Fee) + Platform_Minimum

- (AI_Cost + Human_Talk_Cost + Telecom_Cost)

- (Setup_and_Training + Account_Ops)

The key lever: maximizing `Visits x Success_Fee` while keeping `Human_Talk_Cost` minimal by using AI for the first 2–3 touches per lead.

2.4 Segment KPIs – Real Estate

KPI	Target	Notes
AI Coverage (% of calls)	>50–60%	First-touch qualification
Visit Conversion Uplift	+20–40% vs baseline	Core value proposition
Cost-per-Visit Reduction	20–30%	Includes SquadStack fees

KPI	Target	Notes
Time-to-First-Contact	< 120 seconds	Lead freshness critical
Gross Margin (campaign)	70–80%	After variable costs
Sales Cycle	2–4 months	Developer procurement process

2.5 Edge Cases & Risk Mitigation

Risk	Probability	Mitigation
Developer passes junk leads	High	Contractual valid-lead definitions; feedback loops
CRM integration friction	Medium	Start with CSV/webhook; deep CRM post-pilot
AI distrust on high-ticket call	Medium	AI introduces itself; early human handoff
Long payment cycles	Medium	Monthly minimums provide cash flow baseline
Competitive AI players entering	Medium	Vertical-specific playbooks as moat

2.6 Pilot Plan – Real Estate

Pilot 1 (Weeks 1–12): Proof of Reach and Visit Uplift

Scope:

- 1–2 flagship projects for a single developer (e.g., NCR or Bangalore).
- Lead ingestion via CSV or webhook; no deep CRM integration yet.
- AI handles first-touch within 60–120 seconds of lead arrival.

Success Metrics:

- Reachability: >= 1.5x baseline connect rate.
- Visit conversion: >= 1.2–1.4x baseline visits per 100 leads.
- Cost-per-visit: >= 20% reduction vs developer's own baseline.

Go/No-Go:

- Go: All 3 metrics hit targets → proceed to Pilot 2.
- Yellow: 2 of 3 hit → extend pilot scope and timeline.
- No-Go: <2 metrics → revisit assumptions, pause Real Estate.

Pilot 2 (Weeks 13–24): Portfolio Rollout

Scope: Expand to 2–3 projects or a second developer logo.

Metrics:

- Signed ACV >= ₹40–60L range.
- Repeatability of visit uplift across different project types.
- Operational load manageable for SquadStack delivery team.

Part 3 – Segment 2: EdTech

3.1 Market Opportunity: TAM → SAM → SOM

TAM (India, EdTech support and sales telecalling)

India's EdTech market is large and telecalling-heavy. Sales and support combined often consume a significant share of operating expenditure at major platforms. Large B2C EdTech platforms in test prep and up-skilling each employ hundreds of agents for combined inbound support and outbound sales.

SAM

Focus: 20–30 platforms with significant user bases (1M+ learners or large hybrid/university programs). Each with potential telecalling spends equivalent to ₹10–30Cr/year.

SOM (Year 1 Obtainable)

- Realistic Y1: 1–3 logos → **₹1–3Cr**
- Ideal Y1: 4–6 logos → **₹4–8Cr**

3.2 Demand Validation – Is the Problem Real?

Core pain points:

- **Support burns revenue.** At large platforms, support can reach 15–25% of revenue — making it the 2nd or 3rd largest cost line.
- **Low outbound conversion.** Telecalling conversion from raw lead to paid enrollment often sits at 2–3%, with heavy agent churn and burnout.
- **Inbound repetitive queries.** A large share of support tickets are standardized — batch timings, pricing, access issues — ideal for automation.

Call-type breakdown (typical B2C EdTech):

Call Type	Share	AI-Suitable?
FAQ (schedule, pricing, access issues)	40–50%	Yes
Billing, refunds	15–20%	Yes
Course counseling and enrollment	20–25%	Human needed
Emotional / motivational support	10–15%	Human needed

Verdict: Strong demand, but AI can only handle ~60–65% of calls. Human counselors remain essential for high-value conversions.

3.3 Indicative Unit Economics – EdTech

Example: Large EdTech platform (1M+ learners)

- Potential ACV per major platform: ₹50–100L depending on scope and volumes.
- AI handles bulk of FAQs and drip follow-ups.
- However, human counselor minutes per final enrollment remain high.
- **Directional margin: 60–75%**, lower than Real Estate because of human-heavy closing.

Key constraint: Unlike Real Estate where AI covers 50%+ of the workload, EdTech AI covers only the support layer — the revenue-critical counseling conversations remain human-led.

3.4 Segment KPIs – EdTech

KPI	Target	Notes
Automation of Tier-1 Queries	60–70%	FAQs, access, scheduling

KPI	Target	Notes
Human Minutes per Enrollment	-20–30% vs baseline	Core unit cost lever
Conversion Uplift	+10–20% on follow-up leads	AI drip re-engagement
CSAT / NPS	4.0+ (no regression)	Brand sensitivity is high
Sales Cycle	2–3 months	Platform procurement

3.5 Edge Cases & Risk Mitigation

Risk	Probability	Mitigation
Parent or student AI distrust	Medium	AI introduces itself; instant human escalation
Curriculum complexity	Medium	Scope only to logistics queries initially
Brand risk from hard-sell bot	High	Strict guardrails; no persuasion in AI layer
Counselor churn drag	High	Reduce dependency with AI triage first

3.6 Pilot Plan – EdTech

Pilot 1 (Weeks 1–8): Support Automation Layer

Scope: 1–2 workflows only (e.g., class access issues, batch schedule queries).

Success Metrics:

- = 60–70% of targeted queries handled fully by AI.
- CSAT maintained at 4.0+.
- 30–40% reduction in human ticket volume for covered topics.

Pilot 2 (Weeks 9–20): Re-engagement and Sales-Adjacent Use Cases

Scope: AI-driven re-engagement of dropped-off leads, warm transfer to human counselors.

Metrics: Incremental enrollment rate uplift; net reduction in human minutes per enrollment.

Part 4 – Segment 3: Healthcare (Clinics and Diagnostics)

4.1 Market Opportunity: TAM → SAM → SOM

TAM (India, clinics and diagnostics telecalling)

Mid-sized clinics and diagnostic chains frequently miss inbound calls during working hours because front-desk staff are occupied. A study across 40+ Indian clinics showed missed-call rates of 18–27%, translating to approximately ₹7.4L/year of lost revenue per clinic. Across thousands of organized clinics and labs in Tier 1/2 cities, this represents hundreds of crores in recoverable value.

SAM

Focus: Organized clinic chains and mid-to-large diagnostic networks using some digital infrastructure — those open to outsourcing first-line call handling.

SOM (Year 1 Obtainable)

- Realistic Y1: 10–20 clinics or 1–2 small chains → **₹1–2Cr**
- Ideal Y1: 3–5 chains → **₹4–6Cr**

4.2 Demand Validation – Is the Problem Real?

Core pain points:

- Consistent missed calls during business hours, causing appointment losses to competitors.
- No dedicated 24/7 call center or AI capability.
- Appointment reminders and no-show follow-ups are entirely manual.

Call-type breakdown (typical clinic):

Call Type	Share	AI-Suitable?
Appointment booking	40–50%	Yes
Rescheduling and cancellations	15–20%	Yes

Call Type	Share	AI-Suitable?
Test/price information	15–20%	Yes
Reminders and no-show follow-up	10–15%	Yes
Medical queries / emergencies	5–10%	Must escalate

Verdict: This is the highest AI-coverage use case. 85–90% of calls are transactional and standardized. The main risk is clinical liability, not complexity.

4.3 Indicative Unit Economics – Healthcare

Example: Clinic chain (10 clinics)

- Each clinic loses ~₹7.4L/year from missed calls → chain-level leakage ~₹74L/year.
- SquadStack pricing: modest base fee per clinic + per-appointment or per-handled-call fee.
- ARPU per chain: ₹15–30L, with high automation and minimal human involvement.
- **Directional margin: 65–80%.**

The constraint is ACV, not margin. To matter at scale, SquadStack needs either large chains or a high-volume, low-touch SMB product offering.

4.4 Segment KPIs – Healthcare

KPI	Target	Notes
Missed-Call Rate Reduction	>= 50%	Core value proof
Appointments per 100 Calls	+20% vs baseline	Revenue recovered
No-Show Rate Reduction	-15–25%	Reminder automation
AI Coverage	>85% of handled calls	Logistics only, no medical
Integration Friction	Medium-High	Varies by clinic tech stack

4.5 Edge Cases & Risk Mitigation

Risk	Probability	Mitigation
No CRM or tech stack	High	Virtual call-forwarding setup, no EMR needed
AI gives medical advice	High risk	Strict guardrails; AI only handles logistics
Low ACV per clinic	High	Focus on chains; minimum 5-clinic accounts
Regulatory sensitivity (HIPAA-equiv)	Medium	Data handling contracts; local storage

4.6 Pilot Plan – Healthcare

Pilot 1 (Weeks 1–8): Missed-Call Recovery

Scope: 5–10 clinics from a single chain. Inbound calls forwarded to SquadStack's AI+agent layer when front desk misses the call.

Success Metrics:

- Share of missed calls recovered: >= 70%.
- Incremental appointments booked per week.
- Zero incidents of AI providing clinical advice.

Pilot 2 (Weeks 9–20): Full First-Line Handling

Scope: Expand to first-line call answering during defined hours.

Metrics: Full call-answer rate, appointment conversion rate, cost per appointment booked.

Part 5 – Comparative Final Analysis

5.1 Head-to-Head Comparison

Metric	Real Estate	EdTech	Healthcare
ARPU / ACV (per logo)	₹40–80L	₹50–100L	₹15–30L
Gross Margin Band	70–80%	60–75%	65–80%
AI Automation Potential	High (top-funnel)	Medium (support layer)	Very High (logistics)
Sales Cycle	2–4 months	2–3 months	1–3 months
Buyer Urgency	High	High	Medium-High
Integration Friction	Medium	Low–Medium	Medium–High
Human Skill Intensity	Medium	High	Low
Outcome Pricing Fit	Very High	Medium	Medium

Strategic conclusion:

Real Estate is the primary beachhead because it uniquely combines high ACV, strong AI leverage at top of funnel, and an outcomes-based pricing structure (per site visit) that is clear, verifiable, and creates maximum margin upside for SquadStack.

Healthcare is an ideal parallel track for a more standardized, AI-heavy product. Lower ACV per account is offset by clean AI coverage and low human dependency. It may eventually become a volume-driven, near-autonomous product.

EdTech has the most attractive deal sizes and the sharpest pain, but SquadStack should enter carefully — restricting scope to support automation first and building the human counselor supply model before committing to full sales process outsourcing.

5.2 Directional Y1–Y3 Revenue Trajectories

Realistic Case (Conservative Execution):

Year	Real Estate	Healthcare	EdTech	Combined
Y1	₹2–4Cr (3–5 logos)	₹1–2Cr (2 chains)	₹1–3Cr (1–3 logos)	₹4–9Cr
Y2	₹8–15Cr	₹3–6Cr	₹3–7Cr	₹14–28Cr
Y3	₹20–35Cr	₹6–12Cr	₹6–14Cr	₹32–61Cr

Ideal Case (Aggressive Execution + Market Tailwinds):

Year	Real Estate	Healthcare	EdTech	Combined
Y1	₹6–10Cr	₹4–6Cr	₹4–8Cr	₹14–24Cr
Y2	₹20–40Cr	₹8–15Cr	₹10–18Cr	₹38–73Cr
Y3	₹50–80Cr	₹15–25Cr	₹20–35Cr	₹85–140Cr

Part 6 – Risk Assessment

6.1 Key Risks and Mitigations

Risk 1: Automation % lower than expected

- Probability: Medium (20–30%)
- Impact: Margins lower than modeled; human cost grows faster than revenue
- Mitigation: Narrow Pilot 1 scope to clearly standardized workflows only
- Contingency: If <50% automation in any vertical by Week 8, kill that pilot and de-prioritize

Risk 2: Lead or call quality deteriorates under outcome-based pricing

- Probability: High in Real Estate and EdTech
- Impact: SquadStack bears full variable cost but earns zero success fee on junk leads
- Mitigation: Contract definitions for valid lead; daily lead-quality dashboards
- Contingency: Switch to hybrid model (minimum + outcome) if pure outcome shows losses

Risk 3: Integration drag delays pilots

- Probability: High in Healthcare, Medium in Real Estate
- Impact: Time-to-revenue stretched; pilot data arrives late for decision gates
- Mitigation: Use webhook or CSV-based integration for Pilot 1; deep CRM only post-success
- Contingency: Accept manual data entry for first 4 weeks if necessary

Risk 4: Competitor AI contact centers enter these verticals

- Probability: Medium-High (30%)
- Impact: Pricing pressure, need for faster category lock-in
- Mitigation: Build vertical-specific playbooks and QA IP, not generic voice bots
- Contingency: Accelerate case study publication and reference accounts to establish credibility

Risk 5: LLM and Telecom cost inflation

- Probability: Medium (25%)
- Impact: Margin compression on high-volume, low-ACV accounts (Healthcare)
- Mitigation: Ongoing model optimization; fallback to cheaper LLM alternatives
- Contingency: Re-price accounts with pass-through clauses on cost spikes

6.2 Probability-Weighted Risk Summary

Risk	Probability	Impact	Overall Risk
Automation below target	Medium	High	Medium
Lead quality deterioration	High	Medium	Medium
Integration friction	High	Medium	Medium
Competitor entry	Medium	High	Medium
Cost inflation	Medium	Medium	Low-Medium

Overall risk profile: Medium. Manageable with disciplined pilot design and clear Go/No-Go gates.

Part 7 – Go/No-Go Framework and Founder's Office Execution

7.1 Master Timeline and Decision Gates

Phase	Timeline	Action	Decision
Pilot 1	Weeks 1–12	Validate automation %, contact rate, early lift	Go/No-Go #1
Pilot 2	Weeks 13–24	Validate pricing, CAC, LTV/CAC for top verticals	Go/No-Go #2
Scale	Weeks 25+	Scale winners; park or defer remaining verticals	Ongoing review

7.2 Decision Gate 1 (Week 12)

Vertical	Go Criteria	No-Go Criteria
Real Estate	Reachability >= 1.5x; Visit uplift >= 20%	Both below target after full scope
Healthcare	Missed-call recovery >= 70%; zero medical incidents	Recovery <50% or compliance issues
EdTech	Automation of support queries >= 60–65%	<50% automation; CSAT drops below 3.8

7.3 Decision Gate 2 (Week 24)

Vertical	Go Criteria	No-Go Criteria
Real Estate	2+ logos sign at ₹40L+ ACV	<1 logo signs; CAC > ₹3L
Healthcare	1 chain signs at ₹15L+ with renewal intent	CAC > ₹2L; ACV below floor
EdTech	1 platform signs for support scope at ₹50L+	<1 signs; counselor cost unchanged

7.4 Founder's Office Execution: Your 30-Day Action Plan

If brought into the Founder's Office, the translation of this framework into work looks like this:

Days 0–7: Discovery and Validation

- Walk founders and the sales team through the scoring rubric. Stress-test the vertical rankings with internal deal data.
- Collect baselines from existing verticals (fintech, logistics) to benchmark against new segment economics.

Days 7–14: Target Account Selection

- Build a list of 20–30 target accounts across Real Estate, Healthcare, and EdTech using LinkedIn, company databases, and SquadStack's inbound signals.
- Prioritize accounts showing public signals of high digital ad spend (Meta Ads Library, Google Trends).

Days 14–30: Pilot Design and First Conversations

- Draft pilot offer templates for each vertical with explicit pricing, success metrics, and scope.
- Participate in or prepare founder-led first conversations using this framework as supporting material.
- Use lightweight automation (webhooks, Google Sheets pipelines) to ensure Pilot 1 can run without waiting on engineering sprints.

The core value statement for the Founder's Office:

> "I will not only help decide where to go next but build the systems and structure to learn fast, de-risk expansion bets with real pilot data, and translate strategy into ground-level execution."